

GRPO Policy Reward Convergence (AST + CE Executability) Policy Gradient Loss Dynamics under $\varepsilon = 0.2$ Clipping

